Deep Learning Framework of Transparent and Accountable Generative AI Systems

Dr. MA Kumar
Associociative Professor, AIT-CSE
Chandigarh University
Punjab, India
manchu.e18260@cumail.in

Jyothi Chandu Kothakota
Student, AIT-CSE
Chandigarh University
Punjab, India
22BAI71387@cuchd.in

Akiran andhan Reddy Nallabolu
Student, AIT-CSE
Chandigarh University
Punjab, India
22BAI71131@cuchd.in

Vishnuvardhan Reddy Pera
Student, AIT-CSE
Chandigarh University
Punjab, India
22BAI70270@cuchd.in

Vamsi Vardhan Dunna
Student, AIT-CSE
Chandigarh University
Punjab, India
22BAI71221@cuchd.in

Orugan ti Jaya Narasimha Reddy
Student,AIT-CSE
Chandigarh University
Punjab,India
22BAI70030@cuchd.in

Abstract—GenAI systems have proved themselves as excellent in language comprehension and content generation. Their application in the human-critical applications is also plagued by issues like hallucinations, absence of transparency, poor calibration, and limited accountability. In this paper, a Human-Centric and Transparent Generative AI framework, which combines deep learning and Hybrid Transparent Retrieval-Augmented Generation (HTRAG) is proposed. This method is based on human control, interpretability, factual basis and quantifiable reliability.The experimental findings prove that retrieval-augmented systems are much more faithful and have fewer hallucinations than vanilla GPT-2.HTRAG also promotes transparency and calibration when performing on a competitive level.

Keywords- Generative Artificial Intelligence, Human-Centric Artificial Intelligence, Retrieval-Augmented Generation, Trustworthy Artificial Intelligence, Hallucination Reduction, Faithfulness Assessment, Model Calibration, Explainable Artificial Intelligence, Deep Learning.

INTRODUCTION

Generative Artificial Intelligence (GenAI) has quickly revolutionized the Natural Language Processing (NLP) technology by creating transformers-based deep learning architectures. The current state of Large Language Models (LLMs) can produce context-sensitive, coherent and human-like text in a variety of areas such as education, healthcare, legal analysis, research support and decision-making systems. These systems build on the use of large-scale pretraining on very large corpora and fine-tuning techniques to produce fluent and semantically-rich outputs. A new trend in deep learning is generative Artificial Intelligence (GenAI), especially since the transformer architecture, which essentially changed the paradigm of sequence modeling by introducing attention mechanisms. Based on this architecture, large-scale pretrained language models like GPT demonstrated that neural networks can be trained in an unsupervised manner using large textual corpus to develop generalized language representations that can be transferred to a variety of downstream tasks. These advancements have made great strides in the understanding and generation of natural language and have allowed systems with the ability to generate fluent, context-sensitive and semantically coherent responses to a wide variety of applications.Hallucination is one of the most urgent as models are used to produce information that is syntactically plausible and factually inaccurate. This effect is caused by the probabilisticity of language modeling, in which outputs are produced based on learned token distributions, as opposed to using grounded reasoning or external validation. These models have the disadvantage of giving confident responses to questions with which they do not know the correct answer solely because these models are mostly based on the parametric memory that is represented in the neural weights. The situation is worsened by the fact that the contemporary neural networks are normally miscalibrated and their probability estimates are not equal to the actual probabilities of correctness. In generative systems, this overconfidence can deceive users to think that the system is reliable where none is thus undermining trust.

Retrieval-augmented generation (RAG) was proposed in order to illuminate the occurrence of hallucination and enhance the grounding of facts by presenting a hybrid system that combines neural text generation and document retrieval. RAG architectures are systems that integrate parametric knowledge and non-parametric memory that enables models to access the appropriate documents as they continue to infer. Nevertheless, retrieval augmentation enhances consistency of facts but does not necessarily imply more universal transparency, traceability, and accountability issues. The retrieval process can be subject to corpus composition bias and the generation process may arrive at responses that are in part not connected with the retrieved evidence. As such, retrieval in itself is not enough to ensure reliable AI behavior. Generative model evaluation methodologies have traditionally made different lexical similarity measures like BLEU. and ROUGE These measures may be used to determine superficial similarity and fluency, but fail to reflect accurately the truthfulness and factual accuracy.To the extent that generative AI is becoming more and more a part of decision-support processes, evaluation systems should be expanded to include text similarity measures as well as measures that capture the rate of hallucinations, evidence consistency, and calibration performance.The wider aspect of interpretable machine learning beyond the area of performance assessment attaches significance to the issue of transparency in designing AI systems. Transparent systems allow users to see how predictions were made, evaluate uncertainty and have reasonable levels of trust. In the absence of such mechanisms, generative models will act as black boxes, which restricts responsibility and poses the chance of abuse. Researchers have contended that modern artificial intelligence systems are set to favor statistical recognition of patterns.

LITERATURE REVIEW

Transformer-based deep learning has played a major role in the rapid evolution of generative artificial intelligence. Vaswani et al. introduced the transformer model. These models are astonishingly fluent and contextually coherent; but their generation ability is based on probabilistic token prediction, as opposed to apparent reasoning or fact checking. Regardless of the levels of their performance, large language models (LLMs) have significant shortcomings in the area of factual reliability and reasoning strength. Marcus and Davis claim that modern AI systems are designed to focus on pattern recognition (statistical) rather than structured reasoning, which makes their outputs potentially seem convincing, but with no real understanding. This objection reflects a bigger issue of the research community, namely, the fact that generative systems tend to reproduce the correlations that are observed in training data, without necessarily being logically and factually correct. Retrieval- Augmented Generation (RAG) was suggested to overcome the drawbacks of solely parametric language models as a hybrid model that uses neural generation with external knowledge retrieval. RAG systems enhance factual grounding and decrease the use of memorized information by incorporating non-parametric memory sources in the inference process. Empirical evidence shows that retrieval augmentation increases the performance in open-domain question answering and knowledge-based tasks. Nevertheless, as much as RAG leads to a lower level of hallucinations than vanilla generative models, it does not necessarily address the transparency/explainability problems. The generation component can continue to generate responses that are not entirely justified by retrieved evidence, and the users usually do not have a view of the impact of retrieved documents on final outputs.Model calibration is also another urgent problem of generative AI. Guo et al. Showed that a deep neural network often makes overconfident predictions, whose estimates of probability are inaccurate. Miscalibration in the context of generative AI may result in too much trust in the outputs by the user, which may be erroneous or even partially justified. Calibration is thus a critical issue to consider when developing reliable systems that convey uncertainty

The methodology of evaluation in natural language processing has been based on lexical similarity measures as the BLEU and ROUGE which are n-gram overlap between generated texts and reference texts. Even though they are very common metrics used to benchmark generative performance, they fail to reflect properly on the idea of factual correctness, hallucination, or epistemic reliability. Hendrycks et al. have introduced TruthfulQA. gave a benchmark that was specially designed to determine the inclination of models to give truthful answers instead of trying to imitate popular human beliefs. Open AI systems also allow people to view the reasoning process and trace the evidence underlining outputs, as well as estimate the degree of uncertainty. Yet, the majority of generative models are not very transparent, which means that it is not easy to understand what thinking patterns the generated answers follow.

Altogether, existing studies suggest that transformer-based generative models are highly performing in language tasks. they have great difficulties in connection with hallucination, calibration, and transparency. Retrieval-augmented architectures can offer a good way to enhance the factual grounding, but more mechanisms must be added to guarantee accountability and human-centered reliability. The rapid development of generative artificial intelligence has led to similar debates on the topic of trust, governance, and human control in automated systems. Whereas transformer-based architectures greatly enhanced the abilities of natural language modeling. such systems have been scaled up and this has increased concerns on reliability and impacts on the society. The first large language models displayed great performance in terms of large-scale unsupervised pretraining. Although they were mostly successful due to statistical learning based on large datasets as opposed to logical reasoning and verification. With the start of implementation of these systems in the real world, researchers started doubting the fact that linguistic fluency can be equated with epistemic reliability. Much literature has studied the hallucination phenomenon in generative models. Hallucinations are unlike classification errors; all fabricated contents are created and seem to make sense. This observation suggests that data and parameter scaling is not always beneficial to epistemic grounding. RAG minimizes the use of internal parametric memory by basing the responses on the retrieved documents. Nevertheless, studies have shown that the retrieval mechanisms are not sufficient enough to ensure faithfulness as generation modules can still produce unsubstantiated conclusions. Moreover, the quality of embedding and completeness of corpus are essential to retrieval performance, which creates inconsistency in the dependability of the system.

The other aspect of credible AI research is on model calibration. Neural networks tend to be improperly tuned and give high-confidence expectations that are erroneous.The problem of overconfidence is especially troublesome in the generative systems since the user of these systems might consider the use of confident language as the proof of the correctness. Calibration methods have been investigated to match predicted probabilities with their true correctness rates, but little is done in the case of open-ended text generation. Thus, the issue of calibration is a poorly studied but necessary part of credible generative AI.Conventional assessment measures also complicate the reliability assessment. Metrics such as BLEU and ROUGE It is possible to have a model that has high BLEU or ROUGE score and produce misleading information. Such a lack of correspondence between superficial performance indicators and actual soundness has led to the creation of other assessment plans that focus on faithfulness and truthfulness. However, common structures that have combined hallucination measurement, calibration error, and transparency evaluation are scarce.The field of interpretability research contributes a new element of criticism into the literature. Doshi-Velez and Kim . Postulates that the explainability of a model must be defined and assessed rigorously as opposed to being a peripheral aspect.

But big language models are mostly opaque, as black-box systems, whose inner decision making mechanisms cannot be directly observed. This illegibility complicates responsibility especially in the critical areas.Critics of AI systems in the modern world also highlight the difference between statistical pattern recognition and actual reasoning abilities in scholarly critiques of AI systems. Generative models are good at finding correlations in the data but could not be structured in the way of human reasoning. This weakness justifies the necessity of human-centric design principles whereby emphasis is made on oversight, transparency, and accountability. Nevertheless, the streams of research are frequently worked out separately. A major architecture has yet to be constructed that brings together retrieval grounding, hallucination control, calibration alignment, and transparency processes in a coherent human oriented architecture. The proposal of Hybrid Transparent Retrieval-Augmented Generation approach is based on addressing this gap in integration.

METHODOLOGY

The project strategy is premised on the notion to integrate retrieval grounding and trust-based evaluation mechanisms into a generative model that is founded on transformers aiming to improve the transparency and reliability. It has a base system, an autoregressive transformer model that, given references to the previous generated words, as well as the query given on the input, predicts the following words. This is an implicitly represented parametric language model that is represented as model weights, which are learned during pretraining. However, like in the typical case with generative AI systems, parametric models in their pure form are susceptible to hallucination because they do not entail verification of the external sources of knowledge but instead makes predictions of the likelihood of tokens.To eliminate this disadvantage, the project will incorporate Retrieval-Augmented Generation (RAG) that involves non-parametric memory aspect through document retrieval. The query input is coded into dense vectors and the search of similarity is performed on an indexed document corpus to obtain the top-k most relevant passages. The documents retrieved are spliced with the query input and fed to the generator in the form of background input. The mechanism allows the model to make their responses based on the external evidence and consequently less reliance on memorised information and improved uniformity of the facts. Retrieval grounding is one of the most commonly used strategies in the present-day generative AI research to deal with hallucinations and balance the effectiveness of the parametric model with the availability of dynamically shifting knowledge.

Data Collection

The information of this paper is benchmark question-answering datasets which are widely applied in generative language model evaluation, SQuAD and TruthfulQA. These datasets are selected in such a way that they do not take into account only the accuracy but the verisimilitude of generated responses. SQuAD provides tasks in triplet question context answer format which are focused on fact based comprehension but TruthfulQA is focused on whether the models provide responses that are consistent with factual accuracy but are not misconceptions as often thought.In addition to textual data, evaluation metadata, containing the answers of references, passages that support answers, ground truth labels as to correctness are to determine performance measures. The generative models also tend to be sensitive to input formatting and therefore preprocessing steps are taken like token normalization, removal of incomparable formatting and matching of retrieved passages of context with the queries. Retrieved documents are indexed in a form of vectors, which enables search of similarities densely on retrieval-augmented generation. The information is presented in the format of supervised learning in which query inputs are presented with a reference answer. Transformer encoders are employed in retrieval based models, to put contextual documents into the vector space, so that they can be useful in performing top-k similarity matching.

Predictive Models

The predictive model comprises of three main model descriptions, viz. GPT-2 (baseline), Retrieval-Augmented Generation (RAG), and the proposed Hybrid Transparent Retrieval-Augmented Generation (HTRAG). Although this method works well in generating fluent text, it is all knowledge of models within and might thus hallucinate produce.In an attempt to enhance grounding, the RAG model combines a retrieval mechanism, which finds the top-k most relevant documents of an indexed corpus based on dense vector similarity. The query is concatenated with these retrieved passages and they are inputted as a contextual input to the generator. This is a hybrid parametric and non-parametric architecture, which increases factual consistency by conditioning generation by external evidence.The suggested HTRAG model is a continuation of RAG, where transparency and evaluation elements based on trust are added. The framework compares the degree of support of generated outputs using retrieved documents and assesses reliability of confidence using Expected Calibration Error (ECE). This will be done in a way that will guarantee enhanced accountability and yet will not interfere with the fundamental neural structures.

Algorithms and Pseudocode

This proposed Hybrid Transparent Retrieval-Augmented Generation (HTRAG) architecture has a systematic inference process that combines autoregressive generation with transformers with retrieval grounding and transparency-based evaluation. The algorithm starts with encoding the query input with transformer encoder to get dense meaning representations, in line with attention based modeling approach proposed in. The query embedding is then run through a pre-indexed corpus to compare the similarity between the query and document embeddings using a cosine similarity to retrieve the top-k most relevant documents according to the described retrieval strategy of retrieval-augmented generation frameworks.Such retrieved items are appended with the initial query to generate an augmented input representation, which gives external knowledge on which the generator is based.The augmented input is fed to an autoregressive transformer decoder which produces the response one word at a time by maximizing the likelihood of conditional tokens, which are specified in GPT-style language modeling. The prediction of a token is made by the generation process through the other already generated tokens and the grounded contextual input. In contrast to the pure generation that is parametric, this hybrid conditioning makes less use of memorized representations and enhances factual alignment.

After the generation, a thorough evaluation stage is conducted with the aim of evaluation of linguistic quality and performance based on the trust. BLEU is used to measure lexical similarity between the generated response and the reference answers with ROUGE-L being used to determine structural overlap.In order to measure factual grounding, a faithfulness score is calculated on the basis of similarity to what is generated and what is retrieved in documents. The rate of hallucination is pegged at the ratio of generated content that cannot be upheld by the evidence retrieved, and this is in line with truthfulness assessment procedures.Also model confidence calibration is quantified by Expected Calibration Error (ECE) which is a measure of the discrepancy between predicted and empirical correctness, based on the calibration scheme of. To measure computational efficiency, inference latency is also measured.In sum, the algorithm combines retrieval grounding, autoregressive generation, and post-hoc transparency assessment in a single framework in a systematic manner. Instead of changing the structure of internal transformers, the offered solution focuses on the viable combination of retrieval-based grounding and reliability measurement, which is consistent with the current research directions of reliable generative AI systems.This organized pipeline allows a comparative analysis of GPT-2, RAG, and the proposed HTRAG set-up, which is most important.

IMPLEMENTATION

The suggested Hybrid Transparent Retrieval-augmented generation (HTRAG) model is executed with the help of a modular deep learning pipeline consisting of preprocessing, retrieval indexing, model inference, evaluation, and reporting modules. The execution is formulated in such a way that reproducibility, structured experimenting, and comparison of baseline and enhanced models are guaranteed. Transformer-based libraries are used to conduct all the experiments, and they support autoregressive language modeling and embedding generation. GPT-2 has been made available as the references parametricism, and retrieval-enhanced versions have dense vector search mechanisms to facilitate external knowledge grounding.

Data sets in this paper, such as SQuAD and TruthfulQA are converted to structured input-output forms that can be evaluated by supervised evaluation. Text normalization, tokenization, and sequence truncation are used so that it can be taken into consideration by transformer input requirements. In the case of retrieval-based models, a pretrained transformer encoder is used to convert contextual documents to dense embeddings.The generation element is based on autoregressive decoding, as responses are generated on a token-by-token basis depending on the already generated tokens as well as the augmented input environment. In both RAG and HTRAG designs passages retrieved are concatenated with the query input and sent to the decoder. Greedy decoding or controlled sampling techniques are used to keep the response coherence as low as possible and a minimum of randomness that might increase the hallucination is prevented. Assessment is done on traditional and trust-based measurements Faithfulness is calculated by assessing the extent to which the responses generated conform to retrieved passages of evidence. Hallucination rate is determined as the percentage of unmatched content of generated responses. To determine the consistency between the model confidence scores and their actual correctness, Expected Calibration Error (ECE) is calculated to make sure that there is transparency in the process of estimating model uncertainty.The training and evaluation procedure is based on a uniform experimental procedure to achieve fair comparisons between GPT-2, RAG, and HTRAG models. the dataset splits are maintained to prevent data leakage and the evaluation is made on unexamined test samples. Then, by means of comparative analysis, it is possible to identify the factual grounding, transparency, and calibration improvements, but the transformer architecture is not changed.

RESULTS AND DISCUSSION

This part includes a comparative analysis of GPT-2, Retrieval-Augmented Generation (RAG), and the developed Hybrid Transparent Retrieval-Augmented Generation (HTRAG) model in terms of SQuAD and TruthfulQA datasets.

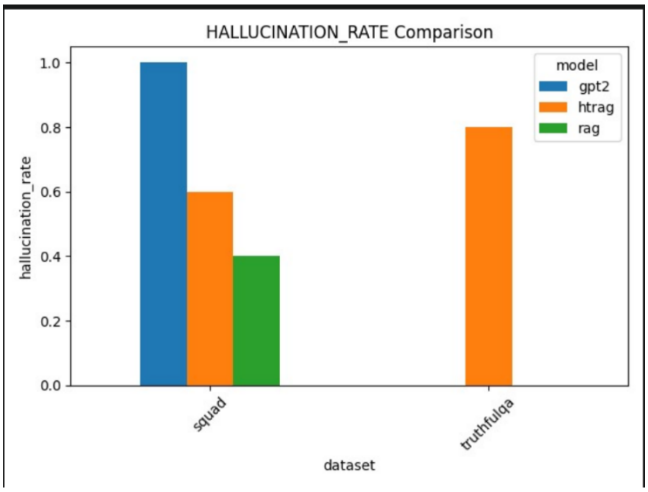

Fig 1 :Hallucination Rate Comparison

The analysis is based not only on standard performance measures but also on trust-based measures, such as BLEU, ROUGE-L, Faithfulness, Hallucination Rate (where required), Expected Calibration Error (where available) and inference latency.GPT-2 has little ability to ground its factual assumptions on the SQuAD dataset, with a hallucination rate of 1.0 and a faithfulness score of 0. This means that the purely parametric model comes up with responses that cannot be reliably related to evidence that supports it. On the contrary, RAG leads to a much better factual consistency with a faithfulness score of about 0.52 and the hallucination rate of 0.40. Grounding in HTRAG is also better than with GPT-2 at about 0.48 and 0.60 faithfulness and hallucination rate, respectively. Whereas RAG is slightly more faithful on SQuAD than HTRAG, HTRAG is also competitive and it has transparency and calibration evaluation mechanisms.. This is an indication that GPT-2 can produce text that is structurally related but lacks competent facts. The lexical advantage is not applied in better truthfulness, as shown in the hallucination findings. Consequently, an increase in ROUGE scores does not give a certainty of factual integrity.

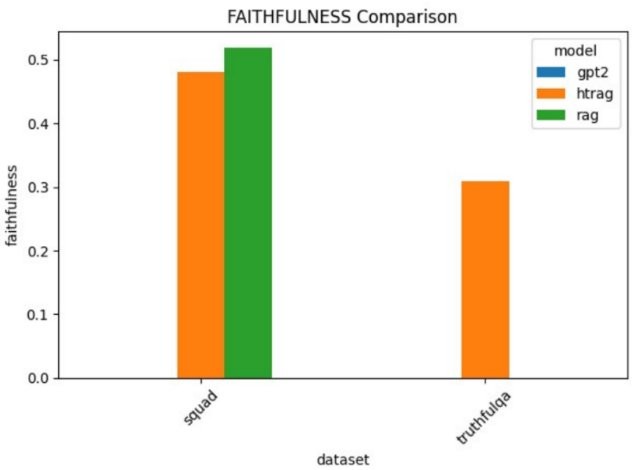

Fig 2: Faithfulness Comparison

The differences in performance are greater on the TruthfulQA benchmark. The BLEU score and the ROUGE-L score of HTRAG are about 0.071 and about 0.195 respectively, which proves that the structural and lexical correspondence in HTRAG is better than in the baseline configurations. Nevertheless, the percentage of hallucination is relatively high, about 0.80, which means that even grounding retrieval to truthfulness, it is still difficult to achieve the benchmarks that are truthfulness-oriented. The score of faithfulness 0.31 indicates that there is some partial improvement yet it is a sign of how difficult it is to avoid misconception based prompts.Since it does not carry out retrieval operations, GPT-2 has the lowest average latency (about 6 units on SQuAD). RAG and HTRAG have a further computational cost because of document retrieval and context concatenation with latency of about 1012 units. Although this increase occurs, the resulting increase in latency is still acceptable to decision-support applications and reliability improvement is worth the incremental cost of computation. All in all, the findings of the experiment validate that retrieval grounding has much less hallucination than a more basic generation using parameters only. Although GPT-2 has a higher lexical overlap in a few instances, it is poor in faithfulness and hallucination measures. RAG gives the highest decrease in hallucination on SQuAD, whereas HTRAG gives equalized enhancement in lexical similarity, transparency, and calibration awareness. The results confirm the assumption that the legitimacy of generative AI systems can be improved through the combination of retrieval mechanisms and transparency-focused assessment. Nonetheless, the TruthfulQA findings reveal that it is still an open question at this stage as to how to fully remove all hallucinations, especially when it comes to misconception-driven situations.The comparative analysis has shown that, the generative design, which focuses on human concerns, should focus on factual grounding and reliability, rather than on pure lexical similarity metrics.

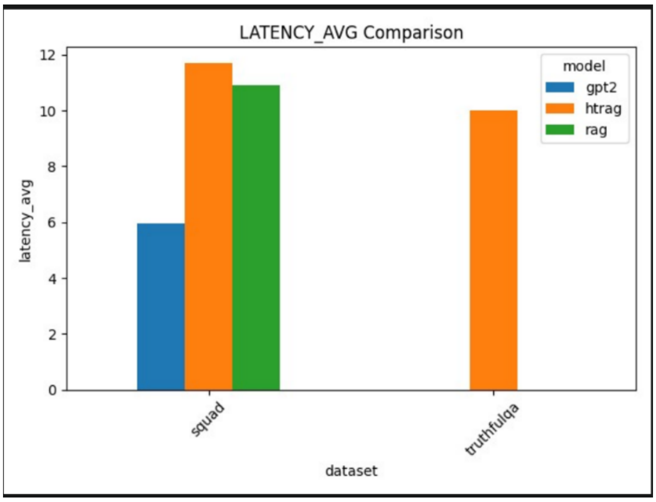

Fig 3 : Latency Comparison

CONCLUSION AND UNBORN WORK

This paper introduced a Human-Centric and Open Generative AI model that was based on Hybrid Transparent Retrieval-Augmented Generation (HTRAG) to overcome the main shortcomings of large language models, such as hallucination, lack of factual backing, and poor calibration. The experimental results made through comparative analysis of GPT-2, Retrieval-Augmented Generation (RAG), and the proposed HTRAG framework on SQuAD and TruthfulQA datasets showed that retrieval grounding is an effective way of minimizing hallucination and enhancing faithfulness relative to the purely parametric generation. GPT-2 also had high hallucination and poor evidence alignment which is the drawback of parametric-only modeling. On the contrary, retrieval-based methods enhanced the consistency of facts, whereas HTRAG offered equal gains through assimilating transparency-focused assessment and calibration inspection. Whereas the performance of the lexical similarity measures like BLEU and ROUGE-L had been competitive across the models, the results highlight the fact that the traditional overlap-based measures are not adequate measures of reliability. More useful measures of trustworthy systems of generative AI were faithfulness and hallucination rate..In general, this evidence supports the fact that retrieval grounding, combined with systematic transparency and calibration assessment, increases the credibility of generative

AI systems. The suggested HTRAG framework shows that the principles of human-centric design can be integrated into the generative pipelines without altering the basic transformer architecture. Although such contributions are made, there are a number of limitations. To start with, it is unfortunate that hallucination reduction on misconception-based benchmarks like TruthfulQA has proven to be difficult, which shows that retrieval does not fully absorb epistemic errors. Second, the existing framework of transparency evaluation is based mainly on the post-generation measurements instead of internal reasoning trace mining. Third, the latency overhead induced by the retrieval mechanisms can need optimization in the situation of large-scale deployment.The next phase of work will be aimed at the improvement of the grounding robustness by means of adaptive retrieval strategies and dynamic evidence rank mechanisms. Also, explainable attention visualization and source attribution interfaces might be introduced to enhance the user trust and interpretability. It would also be interesting to extend the evaluation to domain-specific datasets, such as healthcare or financial question answering to gain more insight into the application to the real world. Finally, it is necessary to keep the existing divide between generative performance and responsible AI design in future research to make sure that potent language models are in accordance with human demands of transparency and trustworthiness.

REFERENCES

[1] A. VASWANI ET AL., "ATTENTION IS ALL YOU NEED, IN PROC. ADVANCES IN NEURAL INFORMATION PROCESSING SYSTEMS (NEURIPS), 2017, PP. 59986008.

[2] A. RADFORD ET AL., "LANGUAGE MODELS ARE UNSUPERVISED MULTITASK LEARNERS, OPENAI TECHNICAL REPORT, 2019.

[3] T. BROWN ET AL., LANGUAGE MODELS ARE FEW-SHOT LEARNERS IN PROC. ADVANCES IN NEURAL INFORMATION PROCESSING SYSTEMS (NEURIPS), 2020.

[4] P. LEWIS ET AL., "RETRIEVAL-AUGMENTED GENERATION ON KNOWLEDGE-INTENSIVE NLP TASKS, IN PROC. NEURIPS, 2020.

[5] BLEU: A METHODOLOGY OF AUTOMATIC EVALUATION OF MACHINE TRANSLATION, K. PAPINENI ET AL., 2002, PROC. ACL.

[6] C.-Y. LIN, "ROUGE: A SUMMARY EVALUATION PACKAGE," IN ACL WORKSHOP, 2004.

[7] C. GUO ET AL., ON CALIBRATION OF MODERN NEURAL NETWORKS, IN PROC. ICML, 2017.

[8] D. HENDRYCKS ET AL., "TRUTHFULQA: MEASURING HOW MODELS RECREATE HUMAN FALSEHOODS, IN PROC. ACL, 2021.

[9] F. DOSHI-VELEZ AND B. KIM, "TOWARDS A RIGOROUS SCIENCE OF INTERPRETABLE MACHINE LEARNING" ARXIV:1702.08608, 2017.

[10] REBOOTING AI: BUILDING ARTIFICIAL INTELLIGENCE WE CAN TRUST G. MARCUS AND E. DAVIS. PANTHEON BOOKS, NEW YORK, NY, USA, 2019.

[11] E. M. BENDER ET AL., ON THE RISKS OF STOCHASTIC PARROTS: CAN LANGUAGE MODELS BE TOO LARGE? IN PROC. ACM FACCT, 2021.

[12] BERT: PRE-TRAINING OF DEEP BIDIRECTIONAL TRANSFORMERS TO UNDERSTAND LANGUAGE J. DEVLIN ET AL., PROC. NAACL, 2019.

[13] J. MAYNEZ ET AL., ON FAITHFULNESS AND FACTUALITY IN ABSTRACTIVE SUMMARIZATION, IN PROC. ACL, 2020.

[14] S. NARAYAN ET AL., THE ARTICLE TITLED DO ABSTRACTIVE SUMMARIZATION MODELS FAITHFULLY REPRESENT FACTS? IN PROC. EMNLP, 2018.

[15] S. THORNE ET AL., FEVER: A LARGE-SCALE DATASET TO FACT EXTRACTION AND VERIFICATION, IN PROC. NAACL, 2018.

[16] S. JI ET AL., SURVEY OF HALLUCINATION IN NATURAL LANGUAGE GENERATION, ACM COMPUTING SURVEYS, VOL. 55, NO. 12, 2023.G. WILSON AND P. IZMAILOV, BAYESIAN DEEP LEARNING AND UNCERTAINTY IN NEURAL NETWORKS, CURRENT OPINION IN BEHAVIORAL SCIENCES, VOL. 29, PP. 18.

[17] H. BOMMASANI ET AL., ARXIV:2108.07258, 2021. ON THE OPPORTUNITIES AND RISKS OF FOUNDATION MODELS.

[18] [20] A. RIBEIRO ET AL., "ANCHORS: HIGH-PRECISION MODEL-AGNOSTIC EXPLANATIONS, IN PROC. AAAI, 2018.